

Non-Linear Modeling

Chapter 7 lab handout

Last modified: 2026-05-28

This handout adapts the Chapter 7 Python lab from the *Introduction to Statistical Learning with Applications in Python* (ISLP) textbook, with permission from the source available at <https://www.statlearning.com>. The original lab is `Ch07-nonlin-lab.Rmd` in the `intro-stat-learning/ISLP_labs` repository.

Students who prefer **R** can work through the corresponding chapter of Emil Hvitfeldt's [ISLR tidymodels labs](#).

In this lab we demonstrate some of the nonlinear models discussed in this chapter. We use the `Wage` data as a running example, and show that many of the complex non-linear fitting procedures discussed can easily be implemented in Python.

1 Setup

As usual, we start with some of our standard imports.

```
import numpy as np, pandas as pd
from matplotlib.pyplot import subplots
import statsmodels.api as sm
from ISLP import load_data
from ISLP.models import (summarize,
                          poly,
                          ModelSpec as MS)
from statsmodels.stats.anova import anova_lm
```

We again collect the new imports needed for this lab. Many of these are developed specifically for the ISLP package.

```
from pygam import (s as s_gam,
                   l as l_gam,
                   f as f_gam,
                   LinearGAM,
                   LogisticGAM)

from ISLP.transforms import (BSpline,
                             NaturalSpline)
```

```
from ISLP.models import bs, ns
from ISLP.pygam import (approx_lam,
                        degrees_of_freedom,
                        plot as plot_gam,
                        anova as anova_gam)
```

2 Polynomial Regression and Step Functions

Exercise 1. Load the Wage data with `load_data('Wage')`. Store the response (wage) as `y` and the predictor (age) as `age`.

Solution

```
Wage = load_data('Wage')
y = Wage['wage']
age = Wage['age']
```

Throughout most of this lab, our response is `Wage['wage']`, which we have stored as `y` above.

Exercise 2. Use `poly()` inside `MS()` to build a 4th-degree polynomial in `age`, fit an OLS model of `y` on this basis, and print the coefficient summary.

Solution

```
poly_age = MS([poly('age', degree=4)]).fit(Wage)
M = sm.OLS(y, poly_age.transform(Wage)).fit()
summarize(M)
```

	coef	std err	t	P> t
intercept	111.7036	0.729	153.283	0.000
poly(age, degree=4)[0]	447.0679	39.915	11.201	0.000
poly(age, degree=4)[1]	-478.3158	39.915	-11.983	0.000
poly(age, degree=4)[2]	125.5217	39.915	3.145	0.002
poly(age, degree=4)[3]	-77.9112	39.915	-1.952	0.051

`poly()` is a *helper* that sets up the transformation; `Poly()` is the workhorse that computes it. The first line above runs the `fit()` method on `Wage`, which stores any parameters needed by `Poly()` on the training data so that subsequent `transform()` calls (on new data, or inside the plotting function below) use the same basis.

Exercise 3. Build a grid `age_grid` of 100 evenly spaced values spanning the observed range of `age`, and wrap it in a single-column `DataFrame` named `age_df`.

Solution

```
age_grid = np.linspace(age.min(),
                        age.max(),
                        100)
age_df = pd.DataFrame({'age': age_grid})
```

Exercise 4. Write a function `plot_wage_fit(age_df, basis, title)` that fits OLS on the supplied basis, evaluates it on `age_df`, and plots the raw data with the fitted curve and its 95% confidence bands.

Solution

```
def plot_wage_fit(age_df,
                  basis,
                  title):


    X = basis.transform(Wage)
    Xnew = basis.transform(age_df)
    M = sm.OLS(y, X).fit()
    preds = M.get_prediction(Xnew)
    bands = preds.conf_int(alpha=0.05)
    fig, ax = subplots(figsize=(8, 8))
    ax.scatter(age,
               y,
               facecolor='gray',
               alpha=0.5)
    for val, ls in zip([preds.predicted_mean,
                       bands[:, 0],
                       bands[:, 1]],
                       ['b', 'r--', 'r--']):
        ax.plot(age_df.values, val, ls, linewidth=3)
    ax.set_title(title, fontsize=20)
    ax.set_xlabel('Age', fontsize=20)
    ax.set_ylabel('Wage', fontsize=20)
    return ax
```

The `alpha` argument to `ax.scatter()` adds transparency that gives a visual indication of density. `zip()` bundles the three lines (fitted mean and two confidence bands) with their colours and line types for the `for` loop.

Exercise 5. Use your helper to plot the degree-4 polynomial fit.

Solution

```
plot_wage_fit(age_df,
              poly_age,
              'Degree-4 Polynomial');
```



Exercise 6. With polynomial regression we must decide on the degree of the polynomial to use. Fit polynomial models in `age` of degrees 1 through 5, and use `anova_lm()` to test which degree is sufficient.

Solution

```
models = [MS([poly('age', degree=d)])  
           for d in range(1, 6)]  
Xs = [model.fit_transform(Wage) for model in models]  
anova_lm([sm.OLS(y, X_).fit()  
          for X_ in Xs])
```

	df_resid	ssr	df_diff	ss_diff	F	Pr(>F)
0	2998.0	5.022216e+06	0.0	NaN	NaN	NaN
1	2997.0	4.793430e+06	1.0	228786.010128	143.593107	2.363850e-32
2	2996.0	4.777674e+06	1.0	15755.693664	9.888756	1.679202e-03
3	2995.0	4.771604e+06	1.0	6070.152124	3.809813	5.104620e-02
4	2994.0	4.770322e+06	1.0	1282.563017	0.804976	3.696820e-01

The `*` unpacks the list of fitted models for `anova_lm()`'s variable argument list. The p-value comparing linear `models[0]` to quadratic `models[1]` is essentially zero, and the cubic-vs-quartic comparison is roughly 5%, while degree 5 is unnecessary ($p = 0.37$). Either a cubic or a quartic polynomial gives a reasonable fit.

Exercise 7. Repeat the ANOVA, but with a linear term in `education` included in each model alongside a polynomial in `age` of degrees 1, 2, 3.

Solution

```
models = [MS(['education', poly('age', degree=d)])
           for d in range(1, 4)]
XEs = [model.fit_transform(Wage)
        for model in models]
anova_lm(*[sm.OLS(y, X_).fit() for X_ in XEs])
```

	df_resid	ssr	df_diff	ss_diff	F	Pr(>F)
0	2997.0	3.902335e+06	0.0	NaN	NaN	NaN
1	2996.0	3.759472e+06	1.0	142862.701185	113.991883	3.838075e-26
2	2995.0	3.753546e+06	1.0	5926.207070	4.728593	2.974318e-02

The ANOVA approach works whether or not the polynomials are orthogonal, provided the models are *nested*. As an alternative, we could choose the polynomial degree by cross-validation (Chapter 5).

Exercise 8. Now consider predicting whether someone earns more than \$250,000 per year. Fit a logistic regression (use `sm.GLM` with a `Binomial` family) of `wage > 250` on the same 4th-degree polynomial basis in `age`, and produce a plot of the predicted probability vs. `age` (with confidence bands and a rug plot).

Solution


```
X = poly_age.transform(Wage)
high_earn = Wage['high_earn'] = y > 250 # shorthand
glm = sm.GLM(y > 250,
             X,
             family=sm.families.Binomial())
B = glm.fit()
summarize(B)
```

	coef	std err	z	P> z
intercept	-4.3012	0.345	-12.457	0.000
poly(age, degree=4)[0]	71.9642	26.133	2.754	0.006
poly(age, degree=4)[1]	-85.7729	35.929	-2.387	0.017
poly(age, degree=4)[2]	34.1626	19.697	1.734	0.083
poly(age, degree=4)[3]	-47.4008	24.105	-1.966	0.049

(The double assignment `high_earn = Wage['high_earn'] = y > 250` both binds the local variable `high_earn` and stores the indicator as a new column on the `Wage` DataFrame. We'll re-use that column in Exercises 29 and 30.)

```
newX = poly_age.transform(age_df)
preds = B.get_prediction(newX)
bands = preds.conf_int(alpha=0.05)

fig, ax = subplots(figsize=(8, 8))
rng = np.random.default_rng(0)
ax.scatter(age +
            0.2 * rng.uniform(size=y.shape[0]),
            np.where(high_earn, 0.198, 0.002),
            fc='gray',
            marker='|')
for val, ls in zip([preds.predicted_mean,
                    bands[:, 0],
                    bands[:, 1]],
                    ['b', 'r--', 'r--']):
    ax.plot(age_df.values, val, ls, linewidth=3)
ax.set_title('Degree-4 Polynomial', fontsize=20)
ax.set_xlabel('Age', fontsize=20)
ax.set_ylim([0, 0.2])
ax.set_ylabel('P(Wage > 250)', fontsize=20);
```



The `age` values above 250 are shown as gray marks at the top of the plot; those below 250 at the bottom. A small amount of jitter prevents identical-age points from overlapping. This is called a *rug plot*.

Exercise 9. Fit a step-function regression: discretise `age` into 4 quantile-based bins with `pd.qcut()`, build dummy columns with `pd.get_dummies()`, and fit OLS.

Solution

```
cut_age = pd.qcut(age, 4)
summarize(sm.OLS(y, pd.get_dummies(cut_age)).fit())
```


	coef	std err	t	P> t
(17.999, 33.75]	94.1584	1.478	63.692	0.0
(33.75, 42.0]	116.6608	1.470	79.385	0.0
(42.0, 51.0]	119.1887	1.416	84.147	0.0
(51.0, 80.0]	116.5717	1.559	74.751	0.0

`pd.qcut()` chose cutpoints at the 25%, 50% and 75% quantiles, so we get four bins. `pd.get_dummies()` (unlike the usual model-matrix convention) keeps *all* category levels; combined with no intercept, each coefficient is the average wage in its bin. For non-quantile cutpoints, use `pd.cut()` instead.

3 Splines

Exercise 10. Use `BSpline()` from `ISLP.transforms` with `internal_knots=[25, 40, 60]` and `intercept=True` to build a cubic B-spline basis for `age`. What shape is the resulting design matrix?

Solution

```
bs_ = BSpline(internal_knots=[25, 40, 60], intercept=True).fit(age)
bs_age = bs_.transform(age)
bs_age.shape
```

(3000, 7)

A cubic B-spline with 3 interior knots yields a 7-column matrix (3 + degree + 1 with `intercept=True`).

Exercise 11. Now fit a cubic spline regression of `wage` on `age` using the helper `bs()` inside `MS()`. Pass `name='bs(age)'` to clean up the column labels in the summary.

Solution

```
bs_age = MS([bs('age',
               internal_knots=[25, 40, 60],
               name='bs(age)')])
Xbs = bs_age.fit_transform(Wage)
M = sm.OLS(y, Xbs).fit()
summarize(M)
```

	coef	std err	t	P> t
intercept	60.4937	9.460	6.394	0.000
bs(age)[0]	3.9805	12.538	0.317	0.751
bs(age)[1]	44.6310	9.626	4.636	0.000
bs(age)[2]	62.8388	10.755	5.843	0.000
bs(age)[3]	55.9908	10.706	5.230	0.000

```
bs(age)[4] 50.6881 14.402 3.520 0.000
bs(age)[5] 16.6061 19.126 0.868 0.385
```

Note there are 6 spline coefficients rather than 7 because `bs()` defaults to `intercept=False` — the basis would normally be 7 columns, but one is dropped to accommodate the model’s overall intercept.

Exercise 12. Instead of specifying knots, ask `BSpline` for `df=6`. Inspect `internal_knots_` to see where the knots ended up.

Solution

```
BSpline(df=6).fit(age).internal_knots_
```

```
array([33.75, 42.    , 51.   ])
```

With `df=6` and cubic splines, three knots are placed at the 25%, 50% and 75% quantiles of `age` — ages 33.75, 42.0, and 51.0.

Exercise 13. B-splines need not be cubic. Fit a degree-0 spline with `df=3` and compare the coefficients to your earlier `qcut()` step-function fit.

Solution

```
bs_age0 = MS([bs('age',
                 df=3,
                 degree=0)]).fit(Wage)
Xbs0 = bs_age0.transform(Wage)
summarize(sm.OLS(y, Xbs0).fit())
```

	coef	std err	t	P> t
intercept	94.1584	1.478	63.687	0.0
bs(age, df=3, degree=0)[0]	22.3490	2.152	10.388	0.0
bs(age, df=3, degree=0)[1]	24.8076	2.044	12.137	0.0
bs(age, df=3, degree=0)[2]	22.7814	2.087	10.917	0.0

Degree-0 B-splines are piecewise constant — equivalent to the step function from Exercise 9. With `df=3` (degree 0), the knots are again at the same three quantiles. The coefficients look different because of the coding: here the intercept is a column of ones and the remaining three coefficients are *increments* relative to the first bin. The sums approximately recover the per-bin means (slight differences are because `qcut()` uses `while` while `bs()` uses `<` at the boundaries).


Exercise 14. Fit a natural cubic spline with 5 degrees of freedom using `ns()`, and plot the fit with `plot_wage_fit()`.

Solution

```
ns_age = MS([ns('age', df=5)]).fit(Wage)
M_ns = sm.OLS(y, ns_age.transform(Wage)).fit()
summarize(M_ns)
```

	coef	std err	t	P> t
intercept	60.4752	4.708	12.844	0.000
ns(age, df=5)[0]	61.5267	4.709	13.065	0.000
ns(age, df=5)[1]	55.6912	5.717	9.741	0.000
ns(age, df=5)[2]	46.8184	4.948	9.463	0.000
ns(age, df=5)[3]	83.2036	11.918	6.982	0.000
ns(age, df=5)[4]	6.8770	9.484	0.725	0.468

```
plot_wage_fit(age_df,  
              ns_age,  
              'Natural spline, df=5');
```



4 Smoothing Splines and GAMs

Exercise 15. A smoothing spline is a special case of a GAM with squared-error loss and a single feature. Fit `LinearGAM(s_gam(0, lam=0.6))` to predict wage from age. (Remember `pygam` expects a 2-D feature matrix.)

Solution

```
X_age = np.asarray(age).reshape((-1, 1))
gam = LinearGAM(s_gam(0, lam=0.6))
gam.fit(X_age, y)
```


```
LinearGAM(callbacks=[Deviance(), Diffs()], fit_intercept=True,
           max_iter=100, scale=None, terms=s(0) + intercept, tol=0.0001,
           verbose=False)
```

`s_gam(0, ...)` says “apply a smoothing spline to column 0 of the feature matrix”. `lam` is the penalty parameter λ . The `-1` inside `.reshape((-1, 1))` tells NumPy to infer that dimension’s size.

Exercise 16. Investigate how the fit changes with the smoothing parameter. Loop `lam` over `np.logspace(-2, 6, 5)` and plot each resulting fit.

Solution


```
fig, ax = subplots(figsize=(8, 8))
ax.scatter(age, y, facecolor='gray', alpha=0.5)
for lam in np.logspace(-2, 6, 5):
    gam = LinearGAM(s_gam(0, lam=lam)).fit(X_age, y)
    ax.plot(age_grid,
            gam.predict(age_grid),
            label='{:.1e}'.format(lam),
            linewidth=3)
ax.set_xlabel('Age', fontsize=20)
ax.set_ylabel('Wage', fontsize=20);
ax.legend(title=r'$\lambda$');
```




Exercise 17. `pygam` can search for an optimal smoothing parameter. Reproduce the plot from Exercise 16, call `.gridsearch(X_age, y)`, and add the resulting fit on top.

Solution

```
fig, ax = subplots(figsize=(8, 8))
ax.scatter(age, y, facecolor='gray', alpha=0.5)
for lam in np.logspace(-2, 6, 5):
    gam = LinearGAM(s_gam(0, lam=lam)).fit(X_age, y)
    ax.plot(age_grid,
            gam.predict(age_grid),
            label='{:.1e}'.format(lam),
            linewidth=3)
gam_opt = gam.gridsearch(X_age, y)
ax.plot(age_grid,
        gam_opt.predict(age_grid),
        label='Grid search',
        linewidth=4)
ax.set_xlabel('Age', fontsize=20)
ax.set_ylabel('Wage', fontsize=20)
ax.legend(title=r'$\lambda$')
fig
```





Exercise 18. It's often more interpretable to specify the *effective degrees of freedom* than λ . Use `approx_lam()` from `ISLP.pygam` to find a λ giving roughly 4 degrees of freedom for the age term, and verify with `degrees_of_freedom()`.

Solution

```
age_term = gam.terms[0]
lam_4 = approx_lam(X_age, age_term, 4)
age_term.lam = lam_4
degrees_of_freedom(X_age, age_term)
```

```
np.float64(4.000000100000307)
```


These degrees of freedom include the unpenalized intercept and linear term of the smoothing spline

— so there are at least 2 df.

Exercise 19. Now loop over df values [1, 3, 4, 8, 15] (passing df + 1 to `approx_lam()` to account for the intercept), refit, and plot each.

Solution

```
fig, ax = subplots(figsize=(8, 8))
ax.scatter(X_age,
           y,
           facecolor='gray',
           alpha=0.3)
for df in [1, 3, 4, 8, 15]:
    lam = approx_lam(X_age, age_term, df + 1)
    age_term.lam = lam
    gam.fit(X_age, y)
    ax.plot(age_grid,
            gam.predict(age_grid),
            label='{ :d}'.format(df),
            linewidth=4)
ax.set_xlabel('Age', fontsize=20)
ax.set_ylabel('Wage', fontsize=20);
ax.legend(title='Degrees of freedom');
```



A value of $df=1$ corresponds to a straight-line fit.

4.1 Additive Models with Several Terms

Exercise 20. Build a GAM “by hand”: use natural splines (`NaturalSpline`) with $df=4$ for age and $df=5$ for year, get dummies for education, stack them all into a model matrix X_{bh} , and fit OLS.

Solution

```
ns_age = NaturalSpline(df=4).fit(age)
ns_year = NaturalSpline(df=5).fit(Wage['year'])
Xs = [ns_age.transform(age),
      ns_year.transform(Wage['year']),
      pd.get_dummies(Wage['education']).values]
X_bh = np.hstack(Xs)
gam_bh = sm.OLS(y, X_bh).fit()
```

We use all columns of the indicator matrix for `education` (no intercept), and stack the three component matrices horizontally to form `X_bh`.

Exercise 21. Construct a partial-dependence plot for `age` from the by-hand GAM: predict on a matrix where all columns except those for `age` are held at their column means, and the `age` columns are swept across `age_grid`.

Solution

```
age_grid = np.linspace(age.min(),
                        age.max(),
                        100)
X_age_bh = X_bh.copy()[1:100]
X_age_bh[:, 0] = X_bh.mean(0)[None, :]
X_age_bh[:, 4] = ns_age.transform(age_grid)
preds = gam_bh.get_prediction(X_age_bh)
bounds_age = preds.conf_int(alpha=0.05)
partial_age = preds.predicted_mean
center = partial_age.mean()
partial_age -= center
bounds_age -= center
fig, ax = subplots(figsize=(8, 8))
ax.plot(age_grid, partial_age, 'b', linewidth=3)
ax.plot(age_grid, bounds_age[:, 0], 'r--', linewidth=3)
ax.plot(age_grid, bounds_age[:, 1], 'r--', linewidth=3)
ax.set_xlabel('Age')
ax.set_ylabel('Effect on wage')
ax.set_title('Partial dependence of age on wage', fontsize=20);
```




The idea is to create a new prediction matrix in which all but the columns for `age` are held constant at the training-data means; the four `age` columns are filled with the natural spline basis evaluated at the 100 values in `age_grid`.

Exercise 22. Do the same for `year`.

Solution

```
year_grid = np.linspace(Wage['year'].min(),
                        Wage['year'].max(),
                        100)
X_year_bh = X_bh.copy()[ :100]
X_year_bh[:, 0] = X_bh.mean(0)[None, :]
X_year_bh[:, 4:9] = ns_year.transform(year_grid)
preds = gam_bh.get_prediction(X_year_bh)
bounds_year = preds.conf_int(alpha=0.05)
partial_year = preds.predicted_mean
center = partial_year.mean()
partial_year -= center
bounds_year -= center
fig, ax = subplots(figsize=(8, 8))
ax.plot(year_grid, partial_year, 'b', linewidth=3)
ax.plot(year_grid, bounds_year[:, 0], 'r--', linewidth=3)
ax.plot(year_grid, bounds_year[:, 1], 'r--', linewidth=3)
ax.set_xlabel('Year')
ax.set_ylabel('Effect on wage')
ax.set_title('Partial dependence of year on wage', fontsize=20);
```

Partial dependence of year on wage




Exercise 23. Now fit the same GAM using *smoothing* splines via `pygam`. Build `LinearGAM(s_gam(0) + s_gam(1, n_splines=7) + f_gam(2, lam=0))`, where `f_gam` is a factor term for education. (Convert education to numeric codes with `.cat.codes`.) Then plot the partial dependence for age using `plot_gam()` from `ISLP.pygam`.

Solution

```
gam_full = LinearGAM(s_gam(0) +  
                     s_gam(1, n_splines=7) +  
                     f_gam(2, lam=0))  
Xgam = np.column_stack([age,  
                        Wage['year'],  
                        Wage['education'].cat.codes])  
gam_full = gam_full.fit(Xgam, y)
```

```
fig, ax = subplots(figsize=(8, 8))  
plot_gam(gam_full, 0, ax=ax)  
ax.set_xlabel('Age')  
ax.set_ylabel('Effect on wage')  
ax.set_title('Partial dependence of age on wage - default lam=0.6', fontsize=20);
```

Partial dependence of age on wage - default lam=0.6



`year` has only 7 unique values, so we use 7 basis functions. `lam=0` on the factor term avoids any

shrinkage. The default `lam=0.6` for `age` produces a noticeably wiggly fit.

Exercise 24. Re-tune `gam_full` so that the `age` and `year` smoothers each have ~4 degrees of freedom (use `approx_lam()` with `df=4+1`), then re-plot the partial dependence for `age` and for `year`.

Solution


```
age_term = gam_full.terms[0]
age_term.lam = approx_lam(Xgam, age_term, df=4 + 1)
year_term = gam_full.terms[1]
year_term.lam = approx_lam(Xgam, year_term, df=4 + 1)
gam_full = gam_full.fit(Xgam, y)
```

Note that updating `age_term.lam` updates `gam_full.terms[0]` in place.

```
fig, ax = subplots(figsize=(8, 8))
plot_gam(gam_full,
         1,
         ax=ax)
ax.set_xlabel('Year')
ax.set_ylabel('Effect on wage')
ax.set_title('Partial dependence of year on wage', fontsize=20)
```

```
Text(0.5, 1.0, 'Partial dependence of year on wage')
```


Partial dependence of year on wage



Exercise 25. Plot the partial dependence for `education` — a categorical term gets a different style of plot, suited to a set of fitted constants per level.

Solution

```
fig, ax = subplots(figsize=(8, 8))
ax = plot_gam(gam_full, 2)
ax.set_xlabel('Education')
ax.set_ylabel('Effect on wage')
ax.set_title('Partial dependence of wage on education',
             fontsize=20);
ax.set_xticklabels(Wage['education'].cat.categories, fontsize=8);
```



4.2 ANOVA Tests for Additive Models

Exercise 26. In all our models, the function of `year` looks fairly linear. Compare three models by ANOVA: a GAM that *excludes* `year` ($\mathcal{M}_1$), a GAM that uses a *linear* term in `year` ($\mathcal{M}_2$), and the full GAM with a smoothing spline in `year` ($\mathcal{M}_3$). What do you conclude?

Solution

```
gam_0 = LinearGAM(age_term + f_gam(2, lam=0))
gam_0.fit(Xgam, y)
gam_linear = LinearGAM(age_term +
                        l_gam(1, lam=0) +
                        f_gam(2, lam=0))
gam_linear.fit(Xgam, y)

anova_gam(gam_0, gam_linear, gam_full)
```

	deviance	df	deviance_diff	df_diff	F	pvalue
0	105402.432206	2991.004005	NaN	NaN	NaN	NaN
1	105134.624170	2990.005190	267.808037	0.998815	7.625329	0.016210
2	105030.684668	2987.007254	103.939502	2.997936	0.986003	0.429637

Note the reuse of `age_term` — that term already has its `lam` set to yield 4 df, and we want to keep that across all three models.

There is compelling evidence that a GAM with a linear `year` term is better than one that omits `year` ($p = 0.002$). However, there is no evidence that a nonlinear function of `year` is needed ($p = 0.435$). So $\mathcal{M}_2$ is preferred.

Exercise 27. Repeat the test for `age`.

Solution

```
gam_0 = LinearGAM(year_term +
                  f_gam(2, lam=0))
gam_linear = LinearGAM(l_gam(0, lam=0) +
                      year_term +
                      f_gam(2, lam=0))
gam_0.fit(Xgam, y)
gam_linear.fit(Xgam, y)
anova_gam(gam_0, gam_linear, gam_full)
```

	deviance	df	deviance_diff	df_diff	F	pvalue
0	109043.809955	2991.000589	NaN	NaN	NaN	NaN
1	107295.111571	2990.000704	1748.698385	0.999884	49.737638	0.000009
2	105030.684668	2987.007254	2264.426903	2.993450	21.513266	0.000026

Very strong evidence that a nonlinear term is required for `age`.

Exercise 28. Print `gam_full.summary()` and use `gam_full.predict()` to generate fitted values on the training set.

Solution

```
gam_full.summary()
```

LinearGAM

```
=====
Distribution:                NormalDist Effective DoF:
Link Function:              IdentityLink Log Likelihood:
14930.261
Number of Samples:          3000 AIC:
                               AICc:
                               GCV:
                               Scale:
                               Pseudo R-Squared:
=====
```

Feature Function	Lambda	Rank	EDoF	P > x	S
s(0)	[465.0491]	20	5.1	1.11e-	
16 ***					
s(1)	[2.1564]	7	4.0	8.10e-	
03 **					
f(2)	[0]	5	4.0	1.11e-	
16 ***					
intercept		1	0.0	1.11e-	
16 ***					

```
=====
Significance codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

WARNING: Fitting splines and a linear function to a feature introduces a model identifiability which can cause p-values to appear significant when they are not.

WARNING: p-values calculated in this manner behave correctly for un-penalized models or models with known smoothing parameters, but when smoothing parameters have been estimated, the p-values are typically lower than they should be, meaning that the tests reject the null too re

```
Yhat = gam_full.predict(Xgam)
```


Exercise 29. Fit a logistic GAM for `high_earn` (i.e. `wage > 250`) using `LogisticGAM(age_term + l_gam(1, lam=0) + f_gam(2, lam=0))`. Plot the partial dependence for education.

Solution

```
gam_logit = LogisticGAM(age_term +
                        l_gam(1, lam=0) +
                        f_gam(2, lam=0))
gam_logit.fit(Xgam, high_earn)

LogisticGAM(callbacks=[Deviance(), Diffs(), Accuracy()],
             fit_intercept=True, max_iter=100,
             terms=s(0) + l(1) + f(2) + intercept, tol=0.0001, verbose=False)

fig, ax = subplots(figsize=(8, 8))
ax = plot_gam(gam_logit, 2)
ax.set_xlabel('Education')
ax.set_ylabel('Effect on wage')
ax.set_title('Partial dependence of wage on education',
             fontsize=20);
ax.set_xticklabels(Wage['education'].cat.categories, fontsize=8);
```



The fit looks flat with very wide error bars on the first category.

Exercise 30. Cross-tabulate `high_earn` against `education`. What do you see?

Solution

```
pd.crosstab(Wage['high_earn'], Wage['education'])
```

	1. < HS Grad	2. HS Grad	3. Some College	4. College Grad	5. Advanced Degree
education					
high_earn					
False	268	966	643	663	381
True	0	5	7	22	45

There are *no* high earners in the lowest education category, so the logistic GAM has nothing to estimate there.

Exercise 31. Refit the logistic GAM excluding observations in the “1. < HS Grad” category, then plot the partial dependence for `education`, `year`, and `age`.

Solution

```
only_hs = Wage['education'] == '1. < HS Grad'
Wage_ = Wage.loc[~only_hs]
Xgam_ = np.column_stack([Wage_['age'],
                          Wage_['year'],
                          Wage_['education'].cat.codes - 1])
high_earn_ = Wage_['high_earn']
```


We subtract 1 from `cat.codes` to work around a `pygam` bug — the relabeling has no effect on the fit.

```
gam_logit_ = LogisticGAM(age_term +
                          year_term +
                          f_gam(2, lam=0))
gam_logit_.fit(Xgam_, high_earn_)
```

```
LogisticGAM(callbacks=[Deviance(), Diffs(), Accuracy()],
             fit_intercept=True, max_iter=100,
             terms=s(0) + s(1) + f(2) + intercept, tol=0.0001, verbose=False)
```


```
fig, ax = subplots(figsize=(8, 8))
ax = plot_gam(gam_logit_, 2)
ax.set_xlabel('Education')
ax.set_ylabel('Effect on wage')
ax.set_title('Partial dependence of high earner status on education', fontsize=20);
ax.set_xticklabels(Wage['education'].cat.categories[1:],
                   fontsize=8);
```


Partial dependence of high earner status on education




```
fig, ax = subplots(figsize=(8, 8))
ax = plot_gam(gam_logit_, 1)
ax.set_xlabel('Year')
ax.set_ylabel('Effect on wage')
ax.set_title('Partial dependence of high earner status on year',
             fontsize=20);
```

Partial dependence of high earner status on year



```
fig, ax = subplots(figsize=(8, 8))
ax = plot_gam(gam_logit_, 0)
ax.set_xlabel('Age')
ax.set_ylabel('Effect on wage')
ax.set_title('Partial dependence of high earner status on age', fontsize=20);
```

Partial dependence of high earner status on age




5 Local Regression

Exercise 32. Finally, use `sm.nonparametric.lowess` to fit local linear regression models with spans of 0.2 and 0.5, and plot them together. Some implementations of GAMs allow local-regression terms; `pygam` does not.

Solution

```
lowess = sm.nonparametric.lowess
fig, ax = subplots(figsize=(8, 8))
ax.scatter(age, y, facecolor='gray', alpha=0.5)
for span in [0.2, 0.5]:
    fitted = lowess(y,
                    age,
                    frac=span,
                    xvals=age_grid)
    ax.plot(age_grid,
            fitted,
            label='{:.1f}'.format(span),
            linewidth=4)
ax.set_xlabel('Age', fontsize=20)
ax.set_ylabel('Wage', fontsize=20);
ax.legend(title='span', fontsize=15);
```



As expected, a span of 0.5 produces a smoother fit than 0.2.